

unemployment-in-india

October 30, 2024

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: data = '/content/drive/MyDrive/Unemployment_Rate_upto_11_2020.csv'
df = pd.read_csv(data)
```

```
[ ]: # Display the first few rows of the dataframe
df.head()
```

```
[ ]:
```

	Region	Date	Frequency	Estimated Unemployment Rate (%)	\
0	Andhra Pradesh	31-01-2020	M	5.48	
1	Andhra Pradesh	29-02-2020	M	5.83	
2	Andhra Pradesh	31-03-2020	M	5.79	
3	Andhra Pradesh	30-04-2020	M	20.51	
4	Andhra Pradesh	31-05-2020	M	17.43	

	Estimated Employed	Estimated Labour Participation Rate (%)	Region.1	\
0	16635535	41.02	South	
1	16545652	40.90	South	
2	15881197	39.18	South	
3	11336911	33.10	South	
4	12988845	36.46	South	

	longitude	latitude
0	15.9129	79.74
1	15.9129	79.74
2	15.9129	79.74
3	15.9129	79.74
4	15.9129	79.74

```
[ ]: '''Step 2: Data Cleaning and Preprocessing Now that we have a better
      ↳ understanding of the dataset, we'll clean and preprocess the data to ensure
      ↳ it's ready for analysis.'''
      # Print and Clean Column Names

      # Print column names to identify any issues
      print("Original Column Names:")
      print(df.columns)

      # Remove any leading or trailing spaces from column names
      df.columns = df.columns.str.strip()

      # Print column names to confirm the changes
      print("\nCleaned Column Names:")
      print(df.columns)
```

Original Column Names:

```
Index(['Region', 'Date', 'Frequency', 'Estimated Unemployment Rate (%)',
      'Estimated Employed', 'Estimated Labour Participation Rate (%)',
      'Region.1', 'longitude', 'latitude'],
      dtype='object')
```

Cleaned Column Names:

```
Index(['Region', 'Date', 'Frequency', 'Estimated Unemployment Rate (%)',
      'Estimated Employed', 'Estimated Labour Participation Rate (%)',
      'Region.1', 'longitude', 'latitude'],
      dtype='object')
```

```
[ ]: # Convert 'Date' to Datetime Format

      # Verify the presence of the 'Date' column
      if 'Date' in df.columns:
          # Remove any leading or trailing spaces from the 'Date' column values
          df['Date'] = df['Date'].str.strip()

          # Convert 'Date' to datetime format
          df['Date'] = pd.to_datetime(df['Date'], format='%d-%m-%Y')

          # Display the data types to confirm the changes
          print("\nData types after conversion:")
          print(df.dtypes)

          # Display the first few rows of the dataframe
          df.head()
      else:
          print("The 'Date' column was not found. Please check the dataset for any
          ↳ discrepancies.")
```

Data types after conversion:

Region	object
Date	datetime64[ns]
Frequency	object
Estimated Unemployment Rate (%)	float64
Estimated Employed	int64
Estimated Labour Participation Rate (%)	float64
Region.1	object
longitude	float64
latitude	float64
dtype:	object

```
[ ]: # Check for Missing Values and Outliers

# Check for missing values in the dataset
missing_values = df.isnull().sum()
print("Missing values in each column:")
print(missing_values)

# Basic statistics to identify any outliers
print("\nSummary statistics:")
print(df.describe())
```

Missing values in each column:

Region	0
Date	0
Frequency	0
Estimated Unemployment Rate (%)	0
Estimated Employed	0
Estimated Labour Participation Rate (%)	0
Region.1	0
longitude	0
latitude	0
dtype:	int64

Summary statistics:

	Date	Estimated Unemployment Rate (%) \
count	267	267.000000
mean	2020-06-16 09:15:30.337078528	12.236929
min	2020-01-31 00:00:00	0.500000
25%	2020-03-31 00:00:00	4.845000
50%	2020-06-30 00:00:00	9.650000
75%	2020-08-31 00:00:00	16.755000
max	2020-10-31 00:00:00	75.850000
std	NaN	10.803283

	Estimated Employed	Estimated Labour Participation Rate (%) \
--	--------------------	---

count	2.670000e+02	267.000000
mean	1.396211e+07	41.681573
min	1.175420e+05	16.770000
25%	2.838930e+06	37.265000
50%	9.732417e+06	40.390000
75%	2.187869e+07	44.055000
max	5.943376e+07	69.690000
std	1.336632e+07	7.845419

	longitude	latitude
count	267.000000	267.000000
mean	22.826048	80.532425
min	10.850500	71.192400
25%	18.112400	76.085600
50%	23.610200	79.019300
75%	27.278400	85.279900
max	33.778200	92.937600
std	6.270731	5.831738

```
[ ]: # Rename Columns (Optional)

# Rename columns for easier reference
df.rename(columns={
    'Estimated Unemployment Rate (%)': 'Unemployment_Rate',
    'Estimated Employed': 'Employed',
    'Estimated Labour Participation Rate (%)': 'Labour_Participation_Rate',
    'Region.1': 'Region_Category'}, inplace=True)
```

```
[ ]: # Display the first few rows after preprocessing
df.head()
```

```
[ ]:
      Region      Date Frequency  Unemployment_Rate  Employed \
0  Andhra Pradesh 2020-01-31         M             5.48  16635535
1  Andhra Pradesh 2020-02-29         M             5.83  16545652
2  Andhra Pradesh 2020-03-31         M             5.79  15881197
3  Andhra Pradesh 2020-04-30         M            20.51  11336911
4  Andhra Pradesh 2020-05-31         M            17.43  12988845


      Labour_Participation_Rate  Region_Category  longitude  latitude
0                41.02                South      15.9129      79.74
1                40.90                South      15.9129      79.74
2                39.18                South      15.9129      79.74
3                33.10                South      15.9129      79.74
4                36.46                South      15.9129      79.74
```

```
[ ]: '''Step 3: Exploratory Data Analysis (EDA) After cleaning and preprocessing the
data, the next step is to perform Exploratory Data Analysis (EDA).'''
```

'''This step will help us understand the patterns, relationships, and key statistics in the data.'''

```
[ ]: # Univariate Analysis
```

```
# Distribution of Unemployment Rate
plt.figure(figsize=(8, 4))
sns.histplot(df['Unemployment_Rate'], bins=30, kde=True)
plt.title('Distribution of Unemployment Rate')
plt.xlabel('Unemployment Rate (%)')
plt.ylabel('Frequency')
plt.show()
```




```
[ ]: # Bivariate Analysis
```

```
# Unemployment Rate vs. Labour Participation Rate
plt.figure(figsize=(8, 4))
sns.scatterplot(x='Labour_Participation_Rate', y='Unemployment_Rate', data=df)
plt.title('Unemployment Rate vs. Labour Participation Rate')
plt.xlabel('Labour Participation Rate (%)')
plt.ylabel('Unemployment Rate (%)')
plt.show()

# Calculate and display the correlation between these variables
correlation = df['Labour_Participation_Rate'].corr(df['Unemployment_Rate'])
```


```
print(f"Correlation between Labour Participation Rate and Unemployment Rate:␣
↪{correlation:.2f}")
```



Correlation between Labour Participation Rate and Unemployment Rate: -0.07

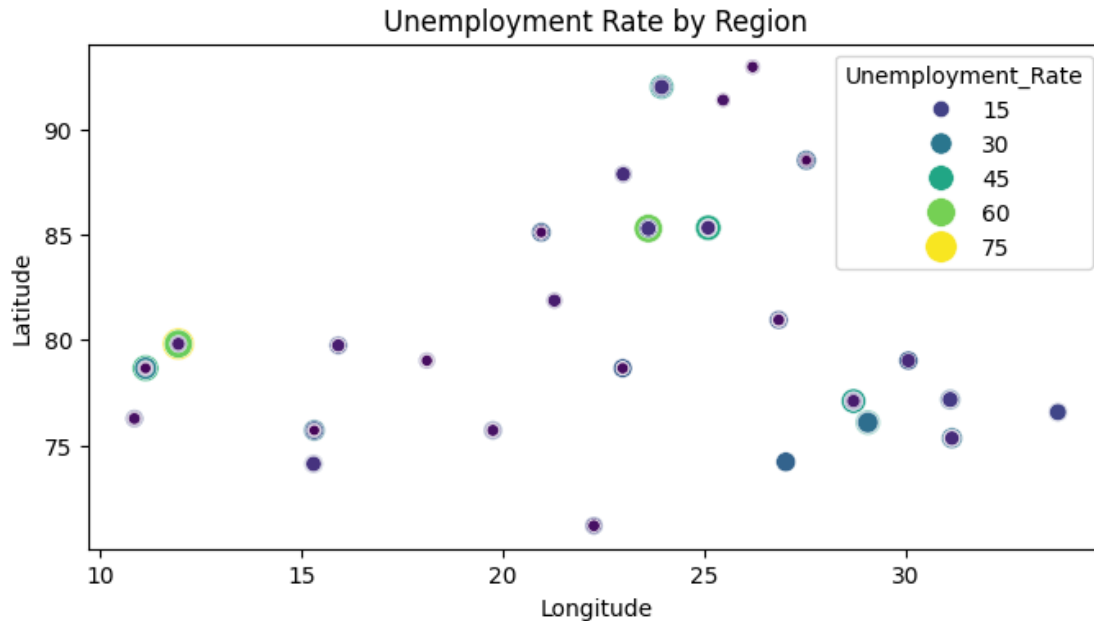
```
[ ]: # Time Series Analysis

# Trend of Unemployment Rate Over Time
plt.figure(figsize=(10,4))
sns.lineplot(x='Date', y='Unemployment_Rate', data=df, marker='o')
plt.title('Unemployment Rate Over Time')
plt.xlabel('Date')
plt.ylabel('Unemployment Rate (%)')
plt.xticks(rotation=45)
plt.show()
```



```
[ ]: # Geographical Analysis

# Unemployment Rate by Region
plt.figure(figsize=(8, 4))
sns.scatterplot(x='longitude', y='latitude', hue='Unemployment_Rate',
               size='Unemployment_Rate', data=df, palette='viridis', sizes=(20, 200))
plt.title('Unemployment Rate by Region')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()
```



```
[ ]: # Hypothesis Testing Once we have explored the data through EDA, we can move on
      ↳ to hypothesis testing to validate any assumptions or insights derived from
      ↳ the data.
```

```
# 1) Formulate Hypotheses Null Hypothesis (H0): There is no significant
      ↳ difference in the unemployment rate across different regions.
```

```
# Alternative Hypothesis (H1): There is a significant difference in the
      ↳ unemployment rate across different regions.
```

```
# 2) Perform the Test
```

```
[ ]: from scipy.stats import f_oneway

# Extracting data for ANOVA
regions = df['Region_Category'].unique()
anova_data = [df[df['Region_Category'] == region]['Unemployment_Rate'] for
      ↳ region in regions]

# Performing the ANOVA test
anova_result = f_oneway(*anova_data)
print(f"ANOVA test result: F-statistic = {anova_result.statistic:.2f}, p-value
      ↳ = {anova_result.pvalue:.4f}")

# Interpretation
if anova_result.pvalue < 0.05:
    print("Result: Reject the null hypothesis. There is a significant
      ↳ difference in unemployment rates across regions.")
```



```

else:
    print("Result: Fail to reject the null hypothesis. No significant
    ↪difference in unemployment rates across regions.")

```

ANOVA test result: F-statistic = 5.04, p-value = 0.0006

Result: Reject the null hypothesis. There is a significant difference in unemployment rates across regions.

```

[ ]: # Model Building (Revised) Now that we have a good understanding of the data
    ↪through our exploratory data analysis (EDA),
    # it's time to build models that can help us predict future unemployment rates.
    # We'll begin by setting up a baseline model and then refine our approach.

```

```

[ ]: # Splitting the Data First, we need to split the data into training and testing
    ↪sets. This will help us evaluate the performance of our model on unseen data.

```

```

[ ]: from sklearn.model_selection import train_test_split

    # Features and target variable
    X = df[['Labour_Participation_Rate', 'Employed', 'longitude', 'latitude']]
    y = df['Unemployment_Rate']

    # Splitting the data
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

    print(f"Training set size: {X_train.shape[0]}")
    print(f"Test set size: {X_test.shape[0]}")

```

Training set size: 213

Test set size: 54

```

[ ]: # Baseline Model: Linear Regression We'll start with a simple Linear Regression
    ↪model as a baseline.
    # This will give us an initial understanding of the relationship between the
    ↪variables.

```

```

[ ]: from sklearn.linear_model import LinearRegression
    from sklearn.metrics import mean_squared_error, r2_score

    # Initialize and train the model
    model = LinearRegression()
    model.fit(X_train, y_train)

    # Make predictions on the test set
    y_pred = model.predict(X_test)

```

```
# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"R-squared (R2): {r2:.2f}")
```

Mean Squared Error (MSE): 88.80
R-squared (R2): 0.06

```
[ ]: # Model Refinement: Feature Engineering If the baseline model's performance is
      ↪not satisfactory, we can refine it by engineering new features or trying
      ↪different models.
      # For instance, we could include interactions between features or try
      ↪polynomial regression
```

```
[ ]: from sklearn.preprocessing import PolynomialFeatures
      from sklearn.pipeline import Pipeline

      # Create a pipeline with polynomial features and linear regression
      poly_model = Pipeline([
          ('poly', PolynomialFeatures(degree=2)),
          ('linear', LinearRegression())
      ])

      # Train the refined model
      poly_model.fit(X_train, y_train)

      # Make predictions
      y_pred_poly = poly_model.predict(X_test)

      # Evaluate the refined model
      mse_poly = mean_squared_error(y_test, y_pred_poly)
      r2_poly = r2_score(y_test, y_pred_poly)

      print(f"Polynomial Model - Mean Squared Error (MSE): {mse_poly:.2f}")
      print(f"Polynomial Model - R-squared (R2): {r2_poly:.2f}")
```

Polynomial Model - Mean Squared Error (MSE): 75.38
Polynomial Model - R-squared (R2): 0.20

```
[ ]: # Model Selection Depending on the results from the baseline and refined
      ↪models, we can choose the best-performing model.
      # If necessary, we might explore other algorithms such as Decision Trees,
      ↪Random Forests, or Gradient Boosting Machines.
```

```
[ ]: # Model Evaluation Finally, we will assess the selected model on the test set,
      ↪ using metrics such as MSE and R-squared.
      # If the model is satisfactory, we can move forward with deploying it or using
      ↪ it for predictive analysis.
```

```
[ ]: # Final model evaluation on the test set
final_model = model # or poly_model, depending on performance

y_final_pred = final_model.predict(X_test)
final_mse = mean_squared_error(y_test, y_final_pred)
final_r2 = r2_score(y_test, y_final_pred)

print(f"Final Model - Mean Squared Error (MSE): {final_mse:.2f}")
print(f"Final Model - R-squared (R2): {final_r2:.2f}")
```

Final Model - Mean Squared Error (MSE): 88.80

Final Model - R-squared (R2): 0.06